

# Assignments 3

NUID : 03137910 Nelson Lefebvre

First of all let's see how does the VGG16 predict the output of a panda image and let's see how the VGG16 is composed.

```
In [5]: import tensorflow as tf
import numpy as np
from tensorflow.keras.layers import Activation, Conv2D, Input, BatchNormalization, Reshape, GlobalAveragePooling
from tensorflow.keras.activations import sigmoid, softmax, relu, tanh
from tensorflow.keras import Sequential
import keras
import tensorflow as tf
tf.compat.v1.disable_eager_execution()
import vis
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import cv2
from keras.preprocessing import image
from keras import backend as K
from keras.applications.vgg16 import preprocess_input, decode_predictions, VGG16, preprocess_input
from vis.utils import utils
from tensorflow.keras.preprocessing.image import load_img

model = VGG16(weights='imagenet')
model.summary()
img = load_img('Panda.jpg', target_size=(224,224))
x = tf.keras.utils.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
print('The most accurate possibility is :', tf.keras.applications.vgg16.decode_predictions(model.predict(x), to
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000

=====  
Total params: 138,357,544  
Trainable params: 138,357,544  
Non-trainable params: 0

C:\ProgramData\Anaconda3\lib\site-packages\keras\engine\training\_v1.py:2356: UserWarning: `Model.state\_updates` will be removed in a future version. This property should not be used in TensorFlow 2.0, as `updates` are applied automatically.

updates=self.state\_updates,  
The most accurate possibility is : [('n02510455', 'giant\_panda', 0.9979412), ('n02447366', 'badger', 0.0011150375), ('n02133161', 'American\_black\_bear', 0.0006159243)]

So we will make a VGG16 "by hand". For that I have retrieved the layers of the VGG16 available with the associate library and I have added different SENet. (I would explain why there are several)

```
In [36]: # Squeeze and Excitation
def se_block(input, channels, r=16):
    # Squeeze
    x = GlobalAveragePooling2D()(input)
    # Excitation
    x = Dense(channels//r, activation="relu")(x)
    x = Dense(channels, activation="sigmoid")(x)
    return tf.keras.layers.Multiply()([input, x])

class SENET_Attn(Layer):
    """
    Channel Attention Block as reported in SENet
    """
    def __init__(self, out_dim, ratio, layer_name="SENET"):
        super(SENET_Attn, self).__init__()
        self.out_dim = out_dim
        self.ratio = ratio
        self.layer_name = layer_name
```

```

def build(self, ratio, layer_name="SENET"):
    self.Global_Average_Pooling = GlobalAveragePooling2D(keepdims= True)
    self.Fully_connected_1_1 = Dense(units= self.out_dim/self.ratio, name=self.layer_name+'_fully_connected',
                                     kernel_initializer="glorot_uniform")

    self.Relu = ReLU()
    self.Fully_connected_2 = Dense(units=self.out_dim, name=layer_name+'_fully_connected2', activation = "tanh")
    self.Sigmoid = Activation("sigmoid")

def call(self, inputs):
    inputs = tf.cast(inputs, dtype = "float32")
    squeeze = self.Global_Average_Pooling(inputs)
    excitation = self.Fully_connected_1_1(squeeze)
    excitation = self.Relu(excitation)
    excitation = self.Fully_connected_2(excitation)
    excitation = self.Sigmoid(excitation)
    excitation = tf.reshape(excitation, [-1,1,1,self.out_dim])

    scale = inputs * excitation
    return scale

Vgg = VGG16(weights='imagenet', include_top=True)

# SENET-RATIO
ratio = 16

input_layer = tf.keras.Input(shape=(224,224,3))
out = Vgg.layers[0](input_layer)
out = Vgg.layers[1](out) # Block 1 of VGG16
out = Vgg.layers[2](out)
out = se_block(out,64)
# out = SENET Attn(out.shape[-1], ratio, )(out) # SENET Attention
#tf.keras.applications.mobilenet.preprocess_input(out)
out = Vgg.layers[3](out)
out = Vgg.layers[4](out) # Block 2 of VGG16
out = Vgg.layers[5](out)
out = se_block(out,128)
# out = SENET Attn(out.shape[-1], ratio, )(out) # SENET Attention
#tf.keras.applications.mobilenet.preprocess_input(out)
out = Vgg.layers[6](out)
out = Vgg.layers[7](out) # Block 3 of VGG16
out = Vgg.layers[8](out)
out = Vgg.layers[9](out)
out = se_block(out,256)
# out = SENET Attn(out.shape[-1], ratio, )(out) # SENET Attention
#tf.keras.applications.mobilenet.preprocess_input(out)
out = Vgg.layers[10](out)
out = Vgg.layers[11](out) # Block 4 of VGG16
out = Vgg.layers[12](out)
out = Vgg.layers[13](out)
out = se_block(out,512)
# out = SENET Attn(out.shape[-1], ratio, )(out) # SENET Attention
#tf.keras.applications.mobilenet.preprocess_input(out)
out = Vgg.layers[14](out)
out = Vgg.layers[15](out) #Block 5 of VGG16
out = Vgg.layers[16](out)
out = Vgg.layers[17](out)
out = se_block(out,512)
# out = SENET Attn(out.shape[-1], ratio, )(out) # SENET Attention
#tf.keras.applications.mobilenet.preprocess_input(out)
out = Vgg.layers[18](out)
out = Vgg.layers[19](out)
out = Vgg.layers[20](out)
out = Vgg.layers[21](out)
out = Vgg.layers[22](out)

model = tf.keras.Model(inputs=input_layer, outputs= out)
model.compile(optimizer = "adam", loss= keras.losses.categorical_crossentropy, metrics = ['accuracy'])
model.summary()

model.fit()

```

Model: "model\_7"

Layer (type)	Output Shape	Param #	Connected to
input_19 (InputLayer)	[(None, 224, 224, 3)]	0	[]
input_18 (InputLayer)	multiple	0	['input_19[0][0]']
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792	['input_18[1][0]']
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928	['block1_conv1[1][0]']
global_average_pooling2d_36 (GlobalAveragePooling2D)	(None, 64)	0	['block1_conv2[1][0]']

dense_72 (Dense)	(None, 4)	260	['global_average_pooling2d_36[0][0]']
dense_73 (Dense)	(None, 64)	320	['dense_72[0][0]']
multiply_36 (Multiply)	(None, 224, 224, 64 )	0	['block1_conv2[1][0]', 'dense_73[0][0]']
block1_pool (MaxPooling2D)	(None, 112, 112, 64 )	0	['multiply_36[0][0]']
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856	['block1_pool[1][0]']
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584	['block2_conv1[1][0]']
global_average_pooling2d_37 (GlobalAveragePooling2D)	(None, 128)	0	['block2_conv2[1][0]']
dense_74 (Dense)	(None, 8)	1032	['global_average_pooling2d_37[0][0]']
dense_75 (Dense)	(None, 128)	1152	['dense_74[0][0]']
multiply_37 (Multiply)	(None, 112, 112, 128)	0	['block2_conv2[1][0]', 'dense_75[0][0]']
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0	['multiply_37[0][0]']
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168	['block2_pool[1][0]']
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080	['block3_conv1[1][0]']
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080	['block3_conv2[1][0]']
global_average_pooling2d_38 (GlobalAveragePooling2D)	(None, 256)	0	['block3_conv3[1][0]']
dense_76 (Dense)	(None, 16)	4112	['global_average_pooling2d_38[0][0]']
dense_77 (Dense)	(None, 256)	4352	['dense_76[0][0]']
multiply_38 (Multiply)	(None, 56, 56, 256)	0	['block3_conv3[1][0]', 'dense_77[0][0]']
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0	['multiply_38[0][0]']
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160	['block3_pool[1][0]']
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808	['block4_conv1[1][0]']
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808	['block4_conv2[1][0]']
global_average_pooling2d_39 (GlobalAveragePooling2D)	(None, 512)	0	['block4_conv3[1][0]']
dense_78 (Dense)	(None, 32)	16416	['global_average_pooling2d_39[0][0]']
dense_79 (Dense)	(None, 512)	16896	['dense_78[0][0]']
multiply_39 (Multiply)	(None, 28, 28, 512)	0	['block4_conv3[1][0]', 'dense_79[0][0]']
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0	['multiply_39[0][0]']
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808	['block4_pool[1][0]']
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808	['block5_conv1[1][0]']
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808	['block5_conv2[1][0]']
global_average_pooling2d_40 (GlobalAveragePooling2D)	(None, 512)	0	['block5_conv3[1][0]']
dense_80 (Dense)	(None, 32)	16416	['global_average_pooling2d_40[0][0]']
dense_81 (Dense)	(None, 512)	16896	['dense_80[0][0]']
multiply_40 (Multiply)	(None, 14, 14, 512)	0	['block5_conv3[1][0]', 'dense_81[0][0]']
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0	['multiply_40[0][0]']
flatten (Flatten)	(None, 25088)	0	['block5_pool[1][0]']

fc1 (Dense)	(None, 4096)	102764544	['flatten[1][0]']
fc2 (Dense)	(None, 4096)	16781312	['fc1[1][0]']
predictions (Dense)	(None, 1000)	4097000	['fc2[1][0]']

```

=====
Total params: 138,435,396
Trainable params: 138,435,396
Non-trainable params: 0

```

```

In [40]: img = load_img('Panda.jpg',target_size=(224,224))
x = tf.keras.utils.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
print('The most accurate possibility is :', tf.keras.applications.vgg16.decode_predictions(model.predict(x), to

The most accurate possibility is : [('n04209239', 'shower_curtain', 0.028491486), ('n03271574', 'electric_fan',
0.018216325), ('n04590129', 'window_shade', 0.016928265)]

```

```

In [26]: img = load_img('Fish.jpg',target_size=(224,224))
x = tf.keras.utils.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
print('The most accurate possibility is :', tf.keras.applications.vgg16.decode_predictions(model.predict(x), to

The most accurate possibility is : [('n03788365', 'mosquito_net', 0.021503624), ('n03495258', 'harp', 0.0212835
9), ('n04589890', 'window_screen', 0.020757237)]

```

```

In [27]: img = load_img('human.jpg',target_size=(224,224))
x = tf.keras.utils.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
print('The most accurate possibility is :', tf.keras.applications.vgg16.decode_predictions(model.predict(x), to

The most accurate possibility is : [('n03788365', 'mosquito_net', 0.08594845), ('n04589890', 'window_screen', 0
.032408703), ('n04590129', 'window_shade', 0.024134178)]

```

```

In [28]: img = load_img('Magpie.jpg',target_size=(224,224))
x = tf.keras.utils.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
print('The most accurate possibility is :', tf.keras.applications.vgg16.decode_predictions(model.predict(x), to

The most accurate possibility is : [('n04589890', 'window_screen', 0.07303518), ('n03788365', 'mosquito_net', 0
.028470566), ('n03271574', 'electric_fan', 0.015849922)]

```

## How is it working ?

SENet (squeeze-and-excitation blocks) is a block who is here to perform dynamic channel-wise feature recalibration. And he is composed of to feature, one normal block which is composed of Gloobal pooling, FC, ReLu, FC, Sigmoid. The we concatenate the Residual block that have (H x W x C size) with the block (who at the end had the 1x1xC size). The goal of using SENet is to increase the rate of prediction but..

## Results

As we can see i implement the SENet with 2 algorithm, both of them seems to work, nevertheless the prediction that i receive is completly inaccurate. I fix the VGG16 but when i add the SENet block the prediction is not accurate. I also tried the CBAM and the STN but without success (same problem, but i add the source is used for implementation).

## Sources :

1 Follow Mr.Zhong advice ) [https://www.tensorflow.org/api\\_docs/python/tf/keras/applications/vgg16/preprocess\\_input](https://www.tensorflow.org/api_docs/python/tf/keras/applications/vgg16/preprocess_input) 2 I asked on tensorflow for the accuracy problem, and i fix half of the problem ) [https://stackoverflow.com/questions/74255600/tensorflow-vgg16-senet-implementation-prediction-problem?noredirect=1#comment131098835\\_74255600](https://stackoverflow.com/questions/74255600/tensorflow-vgg16-senet-implementation-prediction-problem?noredirect=1#comment131098835_74255600) 3 VGG16 prediction ) [https://www.youtube.com/watch?v=YEkriuvrtG0&ab\\_channel=SaptarsiGoswami](https://www.youtube.com/watch?v=YEkriuvrtG0&ab_channel=SaptarsiGoswami) 4 CBAM ) <https://youtu.be/vRYM1KdFtnk> 5 SENet ) <https://youtu.be/MTiqzPdNkFM>

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js