

Assignment 4 :

In [1]:

```
import glob
import imageio
import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
from tensorflow.keras import layers
import time
from tensorflow.keras import layers
from tensorflow.keras import activations
from numpy import expand_dims
from numpy import zeros
from numpy import ones
from numpy import vstack
from numpy.random import randn
from numpy.random import randint
from keras.datasets.mnist import load_data
from keras.optimizers import Adam
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Reshape
from keras.layers import Flatten
from keras.layers import Conv2D
from keras.layers import Conv2DTranspose
from keras.layers import LeakyReLU
from keras.layers import Dropout
from matplotlib import pyplot
from keras.layers import UpSampling2D
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import BatchNormalization
from IPython import display
```

In [2]:

```
(train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1).astype('float32')
train_images = (train_images - 127.5) / 127.5 # Normalize the images to [-1, 1]
BUFFER_SIZE = 60000
BATCH_SIZE = 192
train_dataset = tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

In [3]:

```
def make_generator_model():
    model = Sequential()
    n_nodes = 7 * 7 * 192
    model.add(Dense(n_nodes, input_dim=100))
    model.add(BatchNormalization())
    model.add(layers.Activation(activations.relu))
    model.add(Reshape((7, 7, 192)))
    model.add(Dropout(0.4))
    model.add(UpSampling2D(2))
    model.add(Conv2DTranspose(96, (5,5), strides=1, padding='same'))
    model.add(BatchNormalization())
    model.add(layers.Activation(activations.relu))
    model.add(UpSampling2D(2))
    model.add(Conv2DTranspose(48, (5,5), strides=1, padding='same'))
```

```

model.add(BatchNormalization())
model.add(layers.Activation(activations.relu))
model.add(Conv2DTranspose(24, (5,5), strides=1, padding='same'))
model.add(BatchNormalization())
model.add(layers.Activation(activations.relu))
model.add(Conv2DTranspose(1, (5,5), strides=1, padding='same'))
model.add(Dense(1, activation='sigmoid'))
return model

```

In [4]:

```

generator = make_generator_model()


noise = tf.random.normal([1, 100])
generated_image = generator(noise, training=False)

plt.imshow(generated_image[0, :, :, 0], cmap='gray')

```

Out[4]:

<matplotlib.image.AxesImage at 0x7f3ca17395d0>



In [5]:

```

def make_discriminator_model():
    in_shape=(28,28,1)
    model = tf.keras.Sequential()
    model.add(Conv2D(32, (5,5), strides=(2, 2), padding='same', input_shape=in_shape))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))
    model.add(Conv2D(64, (5,5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))
    model.add(Conv2D(128, (5,5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))
    model.add(Conv2D(256, (5,5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))
    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))
    # compile model
    opt = Adam(lr=0.0002, beta_1=0.5)
    model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
    return model

```

In [6]:

```

discriminator = make_discriminator_model()
decision = discriminator(generated_image)
print (decision)
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)

```

```
total_loss = real_loss + fake_loss
return total_loss
```

```
def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

```
tf.Tensor([[0.49683115]], shape=(1, 1), dtype=float32)
```

```
/usr/local/lib/python3.7/dist-packages/keras/optimizers/optimizer_v2/adam.py:110: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super(Adam, self).__init__(name, **kwargs)
```

In [7]:

```
generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)

checkpoint_dir = './training_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                  discriminator_optimizer=discriminator_optimizer,
                                  generator=generator,
                                  discriminator=discriminator)

EPOCHS = 1
noise_dim = 100
num_examples_to_generate = 16
seed = tf.random.normal([num_examples_to_generate, noise_dim])

@tf.function
def train_step(images):
    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

        generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))

def train(dataset, epochs):
    for epoch in range(epochs):
        start = time.time()

        for image_batch in dataset:
            train_step(image_batch)

        display.clear_output(wait=True)
        generate_and_save_images(generator,
                                  epoch + 1,
                                  seed)

        if (epoch + 1) % 15 == 0:
            checkpoint.save(file_prefix = checkpoint_prefix)

        print ('Time for epoch {} is {} sec'.format(epoch + 1, time.time()-start))

display.clear_output(wait=True)
generate_and_save_images(generator,
                          epochs,
                          seed)
```

```
def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)


    fig = plt.figure(figsize=(4, 4))

    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i+1)
        plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')
        plt.axis('off')

    plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))
    plt.show()
```


In [8]:

```
train(train_dataset, EPOCHS)
```



In [9]:

```
EPOCHS = 10
train(train_dataset, EPOCHS)
```



In [10]:

```
EPOCHS = 20
train(train_dataset, EPOCHS)
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-10-2acff166190e> in <module>
      1 EPOCHS = 20
----> 2 train(train_dataset, EPOCHS)

<ipython-input-7-bab16639e2e4> in train(dataset, epochs)
     42
     43     for image_batch in dataset:
--> 44         train_step(image_batch)
     45
     46         # Produce images for the GIF as you go

/usr/local/lib/python3.7/dist-packages/tensorflow/python/util/traceback_utils.py in error
_handler(*args, **kwargs)
    148     filtered_tb = None
```

```

149     try:
--> 150         return fn(*args, **kwargs)
151     except Exception as e:
152         filtered_tb = _process_traceback_frames(e.__traceback__)

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py in __call__(self, *args, **kwargs)
913
914     with OptionalXlaContext(self._jit_compile):
--> 915         result = self._call(*args, **kwargs)
916
917         new_tracing_count = self.experimental_get_tracing_count()

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py in _call(self, *args, **kwargs)
945         # In this case we have created variables on the first call, so we run the
946         # defunned version which is guaranteed to never create variables.
--> 947         return self._stateless_fn(*args, **kwargs) # pylint: disable=not-callable
948     elif self._stateful_fn is not None:
949         # Release the lock early so that multiple threads can perform the call

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in __call__(self, *args, **kwargs)
2452         filtered_flat_args) = self._maybe_define_function(args, kwargs)
2453         return graph_function._call_flat(
--> 2454             filtered_flat_args, captured_inputs=graph_function.captured_inputs) # pylint: disable=protected-access
2455
2456     @property

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in _call_flat(self, args, captured_inputs, cancellation_manager)
1859         # No tape is watching; skip to running the function.
1860         return self._build_call_outputs(self._inference_function.call(
--> 1861             ctx, args, cancellation_manager=cancellation_manager))
1862         forward_backward = self._select_forward_and_backward_functions(
1863             args,

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py in call(self, ctx, args, cancellation_manager)
500             inputs=args,
501             attrs=attrs,
--> 502             ctx=ctx)
503         else:
504             outputs = execute.execute_with_cancellation(

/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/execute.py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
53         ctx.ensure_initialized()
54         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
--> 55             inputs, attrs, num_outputs)
56     except core._NotOkStatusException as e:
57         if name is not None:

```

KeyboardInterrupt:

In []:

```

EPOCHS = 50
train(train_dataset, EPOCHS)

```

Source : I used the main code in the tensorflow website and i modify it to concord with the template given by the assignment : <https://www.tensorflow.org/tutorials/generative/dcgan>

i also used this template : <https://machinelearningmastery.com/how-to-develop-a-generative-adversarial-network-for-an-mnist-handwritten-digits-from-scratch-in-keras/>

And i found multiple interesting application, used to create different image or other type of media : <https://youtu.be/yz6dNf7X7SA>

https://www.youtube.com/watch?v=uMN85KBV4Rw&ab_channel=mahmoudatifi

<http://prosello.com/papers/text-gans-w17.pdf>

I only do 1 and 10 epochs because it took me a entire day to do the 10 one. So i dont think i will be able to make the 20, 50.. epochs. And the printed image are fully black i dont know exactly why.